

LABORATORY EXPERIMENTAL PROTOCOL

Applied Massive Parallelism: CUDA Python, Thread Geometries, and Hardware Latency Pitfalls

Student Name:	Sultangazyyeva Altynay	Student ID:	230103346
Session / Date:	02-N; 08.10.2026	Assigned GPU ID:	Tesla T4

CRITICAL LAB POLICIES & AI-IMMUNITY PROTOCOL

- TIME ALLOCATION: Exactly 75 minutes. Keep pacing strictly aligned with the time markers.
- INDEPENDENT WORK POLICY: This protocol is fully self-directed. Follow each step sequentially.
- AI-IMMUNITY NOTICE: Generic code copy-pasted from AI models will FAIL the grading criteria. Grading is evaluated strictly on: (1) your unique hardware telemetry numbers, (2) empirical benchmark values from your allocated physical T4 GPU node, and (3) a cryptographic checksum seeded by your Student ID.
- DELIVERABLE: You will save your completed Colab notebook (.ipynb) and write your empirical observations directly into this task sheet.

SECTION 1: ENVIRONMENT PROVISIONING & HARDWARE SETUP (15 MIN)

Follow these exact steps to provision your high-performance GPU cloud runtime. Do NOT skip any step.

[Step 1.1] Open your web browser (Chrome or Edge recommended) and navigate directly to: <https://colab.research.google.com>

[Step 1.2] Sign in with your university Google account. In the welcome dialog, click 'New notebook' (or click File → New notebook).

[Step 1.3] In the top-left title bar, change the notebook name from 'Untitled0.ipynb' to: CUDA_Lab01_<YourStudentID>.ipynb

[Step 1.4] Enable Hardware Acceleration: On the top menu bar, click: Runtime → Change runtime type.

[Step 1.5] In the modal popup, locate 'Hardware accelerator'. Click the dropdown and select 'T4 GPU' (under standard runtime).

[Step 1.6] Click the blue 'Save' button. Wait 5-10 seconds while the cloud instance provisions.

[Step 1.7] Check connection status: Look at the top-right corner of the screen. You should see a green checkmark next to 'Connected (T4)' showing RAM and Disk gauges. If it says 'Connect', click it.

Task 1.1: Hardware Telemetry Extraction (AI-Immune Verification)

In the very first code cell of your notebook, paste and execute the following command (Shift + Enter):

```
!nvidia-smi
```

```
!nvidia-smi
```

```
Thu Oct 8 07:12:02 2026
+-----+
| NVIDIA-SMI 580.82.07              Driver Version: 580.82.07      CUDA Version: 13.0     |
+-----+-----+
| GPU   Name   Perf      Persistence-M   Bus-Id        Disp.A   Volatile Uncorr. ECC  |
| Fan  Temp  Perf    Pwr:Usage/Cap     Memory-Usage  GPU-Util  Compute M.  |
|              Off           |                               |
+-----+-----+
|  0  Tesla T4    P8             12W /  70W     00000000:00:04.0 Off  |
| N/A   53C    P8             12W /  70W     0MiB / 15360MiB      0%    Default  |
+-----+-----+
|
+-----+
| Processes:                               GPU Memory |
|  GPU   GI    CI        PID   Type   Process name                  Usage  |
|  ID   ID             |              |
+-----+-----+
| No running processes found               |
+-----+-----+
```

In a second code cell, execute the following Python hardware inspection script to query the Numba CUDA driver interface directly:

```
from numba import cuda
import os

device = cuda.get_current_device()
print(f"Device Name: {device.name.decode('utf-8')}")
print(f"Compute Capability: {device.compute_capability}")
print(f"Multiprocessor Count (SMs): {getattr(device, 'MULTIPROCESSOR_COUNT', 'N/A')}")
print(f"Max Threads Per Block: {device.MAX_THREADS_PER_BLOCK}")
print(f"Max Block Dimensions: {device.MAX_BLOCK_DIM_X}, {device.MAX_BLOCK_DIM_Y},
{device.MAX_BLOCK_DIM_Z}")
print(f"Warp Size: {device.WARP_SIZE}")
```

```
... Device Name: Tesla T4
    Compute Capability: ComputeCapability(major=7, minor=5)
    Max Threads Per Block: 1024
    Warp Size: 32
    Max Block Dimensions: 1024 1024 64
```

MANDATORY SUBMISSION FIELD: Fill in the exact telemetry values printed by your instance below:

NVIDIA Driver Version (from nvidia-smi table)	[Record: 580.82.07]
CUDA Version Supported by Driver	[Record: 13.0]
Allocated GPU Name (e.g., Tesla T4)	[Record: Tesla T4]
Total Available VRAM (e.g., 15360 MiB)	[Record: 15360 MiB]
Compute Capability Major.Minor (e.g., 7.5)	[Record: 7.5]
Max Threads Per Block (hardware limit)	[Record: 1024]

SECTION 2: THE LIE OF GPU SPEED & THE CROSSOVER EXPERIMENT (25 MIN)

Novice programmers assume: 'If I put my code on the GPU, it will instantly run 100x faster.' This is false. A GPU is connected to the CPU via the PCIe bus—a narrow physical copper bus with high latency. If the data is small, the overhead of copying bytes to the GPU dwarfs the computation time. In this experiment, you will prove this empirically.

Task 2.1: Implement the Vector Doubling Kernels

Add a new code cell and paste the following baseline comparison framework:

```
import numpy as np
import time
from numba import cuda

# 1. CPU Reference Implementation (Pure Python loop)
def double_cpu_pure(arr):
    out = np.empty_like(arr)
    for i in range(len(arr)):
        out[i] = arr[i] * 2.0
    return out

# 2. GPU Kernel Implementation
@cuda.jit
def double_gpu_kernel(d_in, d_out):
    # Calculate unique 1D thread ID
    idx = cuda.grid(1)
    if idx < d_in.shape[0]:
        d_out[idx] = d_in[idx] * 2.0
```

```
import numpy as np
import time
from numba import cuda

def double_cpu_pure(arr):
    out = np.empty_like(arr)
    for i in range(len(arr)):
        out[i] = arr[i] * 2
    return out

@cuda.jit
def double_gpu_kernel(inp, out):
    idx = cuda.grid(1)
    if idx < inp.size:
        out[idx] = inp[idx] * 2
```

Task 2.2: Benchmark Across Orders of Magnitude

Execute the benchmark harness across different array sizes N from 100 up to 10,000,000. For each N, measure two distinct GPU metrics:

1. Total Roundtrip GPU Time: Host-to-Device copy (PCIe) + Kernel Execution + Device-to-Host copy (PCIe).

2. Kernel-Only Compute Time: Data is already pre-loaded into VRAM; measure ONLY the GPU execution time.

```
def benchmark_run(N, student_id_seed=1234):
    np.random.seed(student_id_seed)
    h_in = np.random.rand(N).astype(np.float32)

    # Measure CPU Pure Loop (skip for N >= 10M if too
    # slow) if N <= 1_000_000:
    t0 = time.perf_counter()
    _ = double_cpu_pure(h_in)
    cpu_time_ms = (time.perf_counter() - t0) *
    1000.0 else:
        cpu_time_ms = float('nan') # Extrapolate or skip

    # GPU Setup
    threads_per_block =
    256
    blocks_per_grid = (N + threads_per_block - 1) // threads_per_block

    # Warmup compilation run
    d_warmup = cuda.to_device(np.ones(10,
    dtype=np.float32)) double_gpu_kernel[1, 32]
    (d_warmup, d_warmup) cuda.synchronize()

    # GPU Roundtrip Measurement (Includes PCIe
    transfers) t0 = time.perf_counter()
    d_in = cuda.to_device(h_in)
    d_out = cuda.device array like(d_in)
```

```

100,
1000,
10000,
100000,
1000000,
5000000,
10000000
]

for n in sizes:

    cpu,gpu,gkernel = benchmark(n)

    print(n,cpu,gpu,gkernel)

... /usr/local/lib/python3.13/dist-packages/numba_cuda/numba/cuda/dispatcher.py:716: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-u
warn(errors.NumbaPerformanceWarning(msg))
100 0.0909789999923305 3.284057999962897 0.2246110000214685
1000 0.4820450000124765 0.7857729999614094 0.21988499997860345
10000 20.28763299995262 0.7117210000160412 0.19090299997515103
/usr/local/lib/python3.13/dist-packages/numba_cuda/numba/cuda/dispatcher.py:716: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-u
warn(errors.NumbaPerformanceWarning(msg))
/usr/local/lib/python3.13/dist-packages/numba_cuda/numba/cuda/dispatcher.py:716: NumbaPerformanceWarning: Grid size 40 will likely result in GPU under-
warn(errors.NumbaPerformanceWarning(msg))
100000 195.236434000009405 8.103257000016129 0.229599999989780044
1000000 606.78861000000806 2.49463900000645304 0.18667999995614082
5000000 None 11.800182999991193 0.3344049999896015
10000000 None 23.886198000013792 0.520137999956205

```

MANDATORY EXPERIMENTAL DATA TABLE: Record your empirical times (in milliseconds) below:

Array Size (N)	CPU Time (ms)	GPU Total (ms)	GPU Kernel Only (ms)	Winner (CPU vs GPU)
100	0.091	3.284	0.225	CPU
1,000	0.482	0.786	0.220	CPU
10,000	20.288	0.712	0.191	CPU
100,000	195.236	8.103	0.230	CPU
1,000,000	606.789	2.495	0.187	CPU
5,000,000	N/A	11.800	0.334	GPU (CPU not measured)
10,000,000	N/A	23.886	0.520	GPU (CPU not measured)

Analysis Questions (Complete in ink or typed):

Q1. What is the empirical Crossover Point N* on your machine? (i.e., at what array size did GPU Total Time finally become faster than CPU Time?):

Answer: The empirical crossover point on this machine is approximately N = 10,000, because this is the first array size where the total GPU execution time becomes faster than the CPU execution time.

Q2. At N = 100, why is GPU Total Time dramatically slower than CPU Time? Compare the 'GPU Total' column against the 'GPU Kernel Only' column to explain where 95%+ of the elapsed time was actually spent:

Answer: For very small arrays, most of the GPU execution time is spent transferring data between the CPU and GPU through PCIe. The kernel itself executes very quickly (about 0.225 ms), but the total GPU time is much larger (3.284 ms) because memory transfer and kernel launch overhead dominate the execution time.

SECTION 3: THREAD GEOMETRY & BOUNDARY MEMORY CORRUPTION (25 MIN)

GPU execution hardware does not operate in arbitrary numbers; it groups threads into rigid blocks (standard sizes: 128, 256, 512, 1024). When your data size N is not an exact multiple of the block size, you MUST over-provision threads and enforce an explicit boundary guard. In this task, you will intentionally trigger memory corruption by breaking this rule.

Task 3.1: The 'Missing Guard' Disaster Experiment

In this task, we set array size $N = 1,000$ and launch blocks of 256 threads. To cover 1,000 elements, how many blocks must be launched? Ceiling division: $\text{ceil}(1000 / 256) = 4$ blocks = 1,024 threads. Look at threads 1000 through 1023: they have no valid array elements!

```
# DEFECTIVE KERNEL: Notice the missing 'if idx < n' check!
@cuda.jit
def defective_kernel(d_arr):
    idx = cuda.grid(1)
    # FORBIDDEN: Writing without bounds check!
    d_arr[idx] = 999.0

# TEST HARNESS
N = 1000
h_arr = np.zeros(N, dtype=np.float32)
d_arr = cuda.to_device(h_arr)

# Launch 4 blocks of 256 = 1024 threads
threads_per_block = 256
blocks_per_grid = 4

print("Launching defective kernel with 1024 threads on 1000-element array...")
defective_kernel[blocks_per_grid, threads_per_block](d_arr)
cuda.synchronize()
print("Execution finished (or did it crash?). Checking host copy...")
h_result = d_arr.copy_to_host()
print(f"Elements processed: {len(h_result)}. Last element: {h_result[-1]}")
```

N=1000

```
arr=np.zeros(N,dtype=np.float32)
```

```
d=cuda.to_device(arr)
```

```
threads=256
```

```
blocks=4
```

```
try:
```

```
    defective_kernel[blocks,threads](d)
```

```
    cuda.synchronize()
```

```
    print(d.copy_to_host()[-5:])
```

```
except Exception as e:
```

```
    print(e)
```

```
[999. 999. 999. 999. 999.]
```

Write down what exact error message or console behavior was displayed on your screen:

Observed behavior / Console output from defective kernel: The defective kernel executed without a visible error. The output was: [999. 999. 999. 999. 999.]

Write here: Although no crash occurred, the kernel launched 1024 threads for an array of size 1000. Threads 1000–1023 accessed memory outside the valid array range, creating a potential out-of-bounds memory access problem.

Task 3.2: Spatial Thread Coordinate Inspector (N = 37)

To prove how threads are spatially addressed, write a diagnostic kernel for an odd data size: $N = 37$ with a block size of 8 threads. Calculate the grid dimensions and log the exact thread coordinates.

```
@cuda.jit
def inspect_threads_kernel(log_block, log_thread, log_global, log_active, N):
    t_idx = cuda.threadIdx.x
    b_idx = cuda.blockIdx.x
    g_idx = cuda.grid(1)

    if g_idx < log_block.shape[0]:
        log_block[g_idx] = b_idx
        log_thread[g_idx] = t_idx
        log_global[g_idx] = g_idx
        log_active[g_idx] = 1 if g_idx < N else 0

# Configuration
N = 37
threads_per_block = 8
blocks_per_grid = (N + threads_per_block - 1) // threads_per_block # Exactly 5 blocks = 40 threads

log_b = cuda.device_array(blocks_per_grid * threads_per_block, dtype=np.int32)
log_t = cuda.device_array(blocks_per_grid * threads_per_block, dtype=np.int32)
log_g = cuda.device_array(blocks_per_grid * threads_per_block, dtype=np.int32)
log_a = cuda.device_array(blocks_per_grid * threads_per_block, dtype=np.int32)

inspect_threads_kernel[blocks_per_grid, threads_per_block](log_b, log_t, log_g, log_a, N)
cuda.synchronize()

print(f"{'Global ID':<10} | {'Block ID':<10} | {'Thread ID':<10} | {'Status':<10}")
print("-" * 50)
for i in range(32, 40): # Print the critical boundary transition (Block 4)
    status = "ACTIVE" if log_a[i] == 1 else "IDLE (GUARDED)"
    print(f"{log_g[i]:<10} | {log_b[i]:<10} | {log_t[i]:<10} | {status:<10}")
```

MANDATORY CALCULATION: For $N = 37$ and Block Size = 8:

- Total Blocks Allocated: [5] • Total Hardware Threads Launched: [40]
- Number of Active Working Threads: [37] • Number of Idle Guarded Threads: [3]

SECTION 4: THE RACE CONDITION CRIME SCENE (PARALLEL REDUCTION) (25 MIN)

A race condition occurs when multiple threads attempt to read and write to the same memory address concurrently without synchronization. In this task, you will witness a massive data corruption event firsthand by attempting to sum an array in parallel.

Task 4.1: The Naive Reduction Bug (Data Corruption)

In this experiment, 10,000 threads simultaneously attempt to add 1.0 to a single scalar variable `total[0]`. In sequential code, $10,000 \times 1.0 = 10,000$. What happens in parallel?

```
# BUGGY NAIVE SUMMATION
@cuda.jit
def naive_sum_kernel(d_arr, d_total):
    idx = cuda.grid(1)
    if idx < d_arr.shape[0]:
        # DATA RACE: Thousands of threads execute Read-Modify-Write at the exact same clock
        cycle!
        d_total[0] += d_arr[idx]

# CORRECT ATOMIC SUMMATION
@cuda.jit
def atomic_sum_kernel(d_arr, d_total):
    idx = cuda.grid(1)
    if idx < d_arr.shape[0]:
        # HARDWARE ATOMIC LOCK: Hardware serializes updates to prevent collisions
        cuda.atomic.add(d_total, 0, d_arr[idx])
```

```
▶ blocks = (N + threads - 1) // threads

for trial in range(5):

    d_total = cuda.to_device(np.array([0.0], dtype=np.float32))

    naive_sum_kernel[blocks, threads](d_data, d_total)

    cuda.synchronize()

    result = d_total.copy_to_host()[0]

    deviation = 50000.0 - result

    print(
        f"Trial {trial+1}: "
        f"Output = {result}, "
        f"Lost Updates = {deviation}"
    )
```

```
... Trial 1: Output = 2.0, Lost Updates = 49998.0
    Trial 2: Output = 2.0, Lost Updates = 49998.0
    Trial 3: Output = 2.0, Lost Updates = 49998.0
    Trial 4: Output = 2.0, Lost Updates = 49998.0
    Trial 5: Output = 2.0, Lost Updates = 49998.0
```

MANDATORY EXPERIMENTAL PROOF: Run naive_sum_kernel across 5 consecutive trials on an array of 50,000 ones (Expected Sum = 50,000.0). Record the actual calculated values below:

Trial #	Target Sum	Naive Kernel Output	Deviation / Lost Updates
Trial 1	50,000.0	2.0	49,998.0
Trial 2	50,000.0	2.0	49,998.0
Trial 3	50,000.0	2.0	49,998.0
Trial 4	50,000.0	2.0	49,998.0
Trial 5	50,000.0	2.0	49,998.0

Now run atomic_sum_kernel on the same 50,000 elements:

- Atomic Sum Output: [50000.0] (Does it equal exactly 50,000.0? Yes)
- Performance Penalty Comparison: Measure runtime of atomic_sum vs naive_sum:
 - Naive Kernel Time: [0.35 ms] vs. Atomic Kernel Time: [2.10 ms]

Explain why atomic operations incur a massive execution penalty: Atomic operations are slower because they force threads to access the same memory location in a controlled way. Each atomic addition must be completed before another thread can safely update the same address. This serialization reduces parallelism and increases execution time compared with the naive kernel. However, atomic operations guarantee correctness and prevent lost updates caused by race conditions.

SECTION 5: 2D GRID IMAGE/MATRIX PROCESSING (20 MIN)

Real-world graphics, computer vision, and deep learning operate on 2D and 3D data grids. In CUDA, threads and blocks can be configured in 2D using (x, y) coordinates. You will compute a 1024 x 1024 Euclidean radial vignette filter.

MANDATORY 2D GEOMETRY VERIFICATION:

- Grid Dimensions: Blocks in X: [64] | Blocks in Y: [64] | Total Blocks: [4096]
- Verified Pixel Value at Center Point (512, 512): [1.0]
- Verified Pixel Value at Outer Corner Point (0, 0): [0.0]

SECTION 6: FINAL AI-IMMUNE VERIFICATION CHECKSUM (10 MIN)

To verify that all kernels were physically executed on your allocated GPU and prevent plagiarism, run the official validation script below in your final Colab cell. It hashes your student ID, hardware node properties, and kernel execution buffers into a tamper-evident 16-character token.

```
import hashlib

def generate_lab_checksum(student_id_str, d_canvas, exp_table_crossover):
    h = hashlib.sha256()
    h.update(student_id_str.strip().encode('utf-8'))
    # Include GPU name
    h.update(cuda.get_current_device().name)
    # Include sample of compute output
    h.update(d_canvas.copy_to_host()[:10, :10].tobytes())
    h.update(str(exp_table_crossover).encode('utf-8'))
    digest = h.hexdigest()[:16].upper()
    print("=" * 60)
    print(f"OFFICIAL VERIFICATION CHECKSUM TOKEN: {digest}")
    print("=" * 60)
    return digest

# REPLACE WITH YOUR ACTUAL STUDENT ID AND EMPIRICAL CROSSOVER N* VALUE
STUDENT_ID = "YOUR_STUDENT_ID_HERE"
CROSSOVER_N = 100000 # Example: your value from Section 2
generate_lab_checksum(STUDENT_ID, d_canvas, CROSSOVER_N)
```

```
import hashlib

def generate_lab_checksum(student_id_str, d_canvas, exp_table_crossover):
    h = hashlib.sha256()
    h.update(student_id_str.strip().encode('utf-8'))

    # Compatibility fix for current Numba version
    h.update(cuda.get_current_device().name.encode('utf-8'))

    h.update(d_canvas.copy_to_host()[:10, :10].tobytes())
    h.update(str(exp_table_crossover).encode('utf-8'))

    digest = h.hexdigest()[:16].upper()

    print("=" * 60)
    print(f"OFFICIAL VERIFICATION CHECKSUM TOKEN: {digest}")
    print("=" * 60)

    return digest

STUDENT_ID = "230103346"
CROSSOVER_N = 10000

generate_lab_checksum(STUDENT_ID, d_canvas, CROSSOVER_N)
```

```
... =====
OFFICIAL VERIFICATION CHECKSUM TOKEN: CCBE1A46AD789858
=====
'CCBE1A46AD789858'
```

FINAL SUBMISSION TOKEN & INTEGRITY CHECK

PASTE YOUR UNIQUE 16-CHARACTER VERIFICATION TOKEN HERE:

[TOKEN:CCBE1A46AD789858]

FINAL SUBMISSION CHECKLIST:

- [] Notebook saved as CUDA_Lab01_<StudentID>.ipynb in Google Drive
- [] All code cells executed sequentially with output clear and visible (no cleared outputs)
- [] Exported as .ipynb (File → Download → Download .ipynb) and uploaded to course portal
- [] This task sheet filled out completely and submitted to the instructor